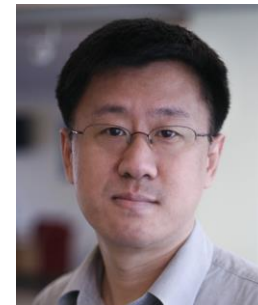


Testing Database Engines via Pivoted Query Synthesis



Manuel Rigger



Zhendong Su

ETH Zurich, Switzerland

11/05/2020



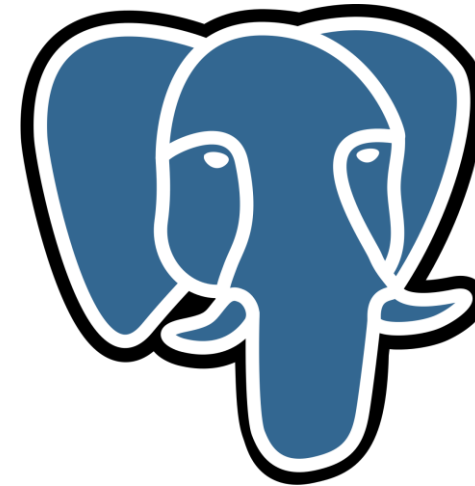
@RiggerManuel @ast_eth

<https://people.inf.ethz.ch/suz/>

Database Management Systems (DBMSs)



PostgreSQL



Database Management Systems (DBMSs)



PostgreSQL



“it is seems likely that there are **over one trillion (1e12) SQLite databases** in active use”

MySQL™

Goal: Find Logic Bugs



Logic bugs: DBMS returns an incorrect result set

Example: SQLite3 Bug

t0

c0
0
1
2
NULL

```
CREATE TABLE t0(c0);  
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;  
INSERT INTO t0 (c0) VALUES (0), (1), (2), (NULL);  
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```

Example: SQLite3 Bug

t0

c0
0
1
2
NULL

```
CREATE TABLE t0(c0);  
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;  
INSERT INTO t0 (c0) VALUES (0), (1), (2), (NULL);  
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```

IS NOT is a “null-safe”
comparison operator

Example: SQLite3 Bug

t0

c0
0
1
2
NULL

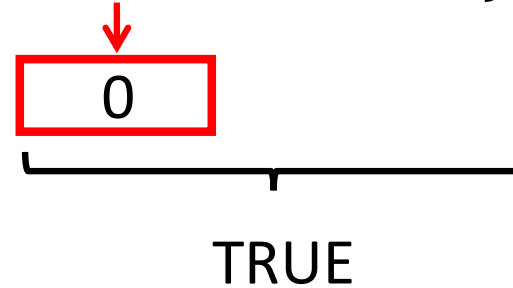
```
CREATE TABLE t0(c0);  
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;  
INSERT INTO t0 (c0) VALUES (0), (1), (2), (NULL);  
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```

Example: SQLite3 Bug

t0

c0
0
1
2
NULL

```
CREATE TABLE t0(c0);  
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;  
INSERT INTO t0 (c0) VALUES (0), (1), (2), (NULL);  
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```



Example: SQLite3 Bug

t0

c0
0
1
2
NULL

```
CREATE TABLE t0(c0);  
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;  
INSERT INTO t0 (c0) VALUES (0), (1), (2), (NULL);  
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```



0

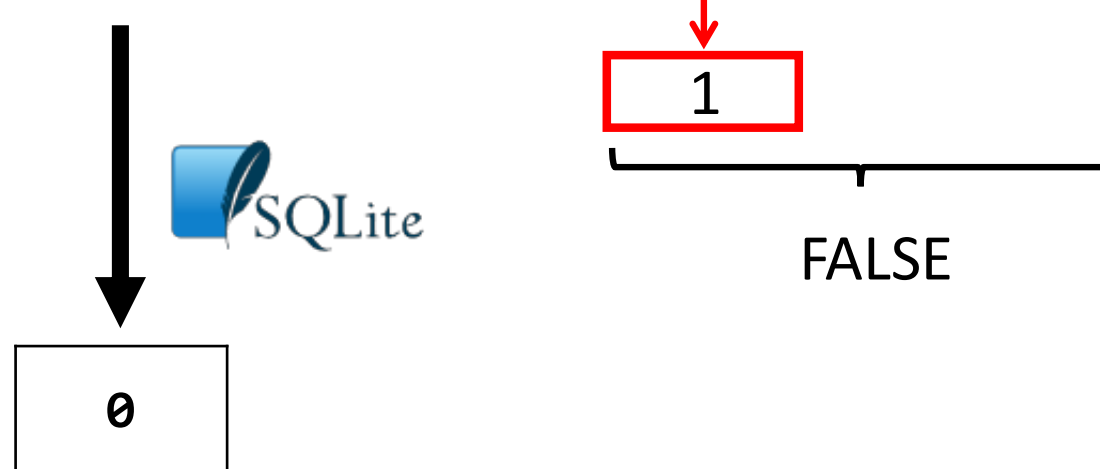
0
TRUE

Example: SQLite3 Bug

t0

c0
0
1
2
NULL

```
CREATE TABLE t0(c0);  
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;  
INSERT INTO t0 (c0) VALUES (0), (1), (2), (NULL);  
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```



Example: SQLite3 Bug

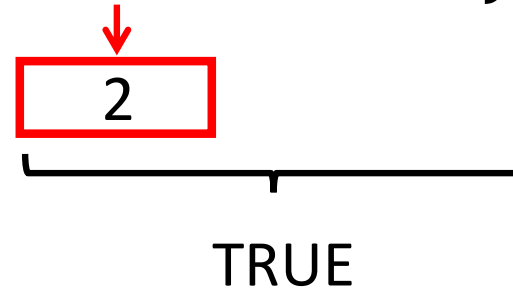
t0

c0
0
1
2
NULL

```
CREATE TABLE t0(c0);  
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;  
INSERT INTO t0 (c0) VALUES (0), (1), (2), (NULL);  
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```



0
2



Example: SQLite3 Bug

t0

c0
0
1
2
NULL

```
CREATE TABLE t0(c0);  
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;  
INSERT INTO t0 (c0) VALUES (0), (1), (2), (NULL);  
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```



↓
NULL
└──────────┘
TRUE


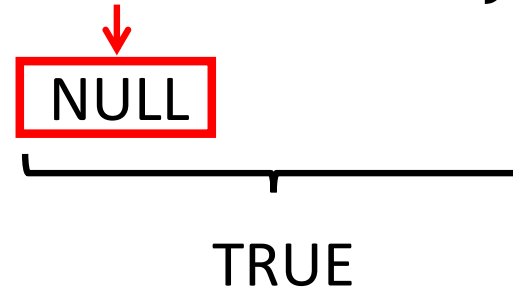
0
2
NULL

Example: SQLite3 Bug

t0

c0
0
1
2
NULL

```
CREATE TABLE t0(c0);  
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;  
INSERT INTO t0 (c0) VALUES (0), (1), (2), (NULL);  
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```



0
2
NULL

Example: SQLite3 Bug

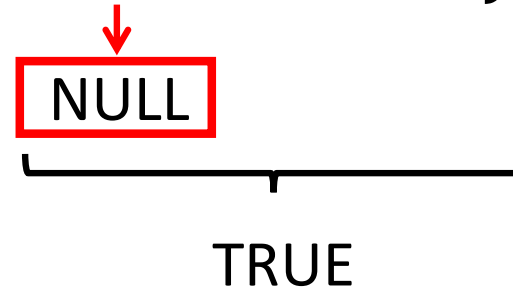
t0

c0
0
1
2
NULL

```
CREATE TABLE t0(c0);  
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;  
INSERT INTO t0 (c0) VALUES (0), (1), (2), (NULL);  
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```



0
2

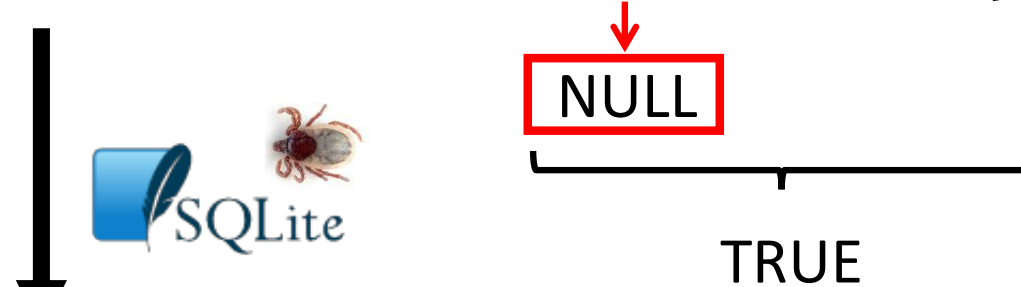


Example: SQLite3 Bug

t0

c0
0
1
2
NULL

```
CREATE TABLE t0(c0);  
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;  
INSERT INTO t0 (c0) VALUES (0), (1), (2), (NULL);  
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```



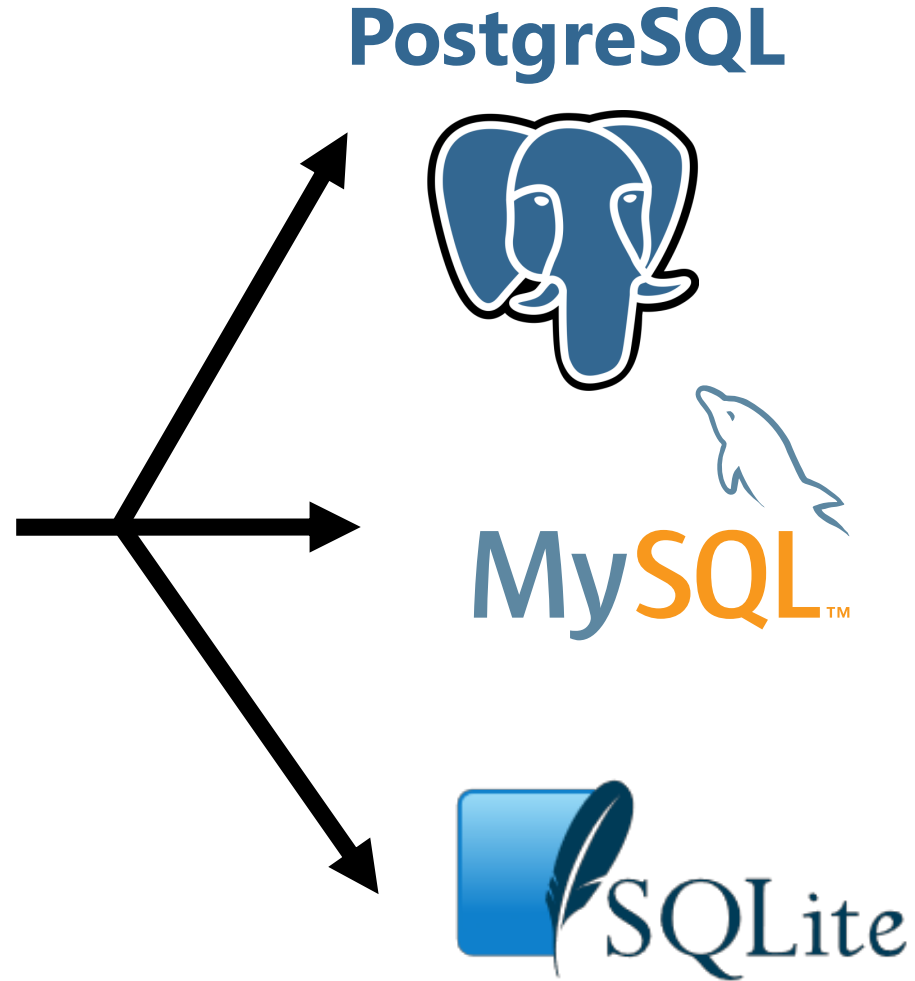
0
2

NULL was not contained
in the result set!



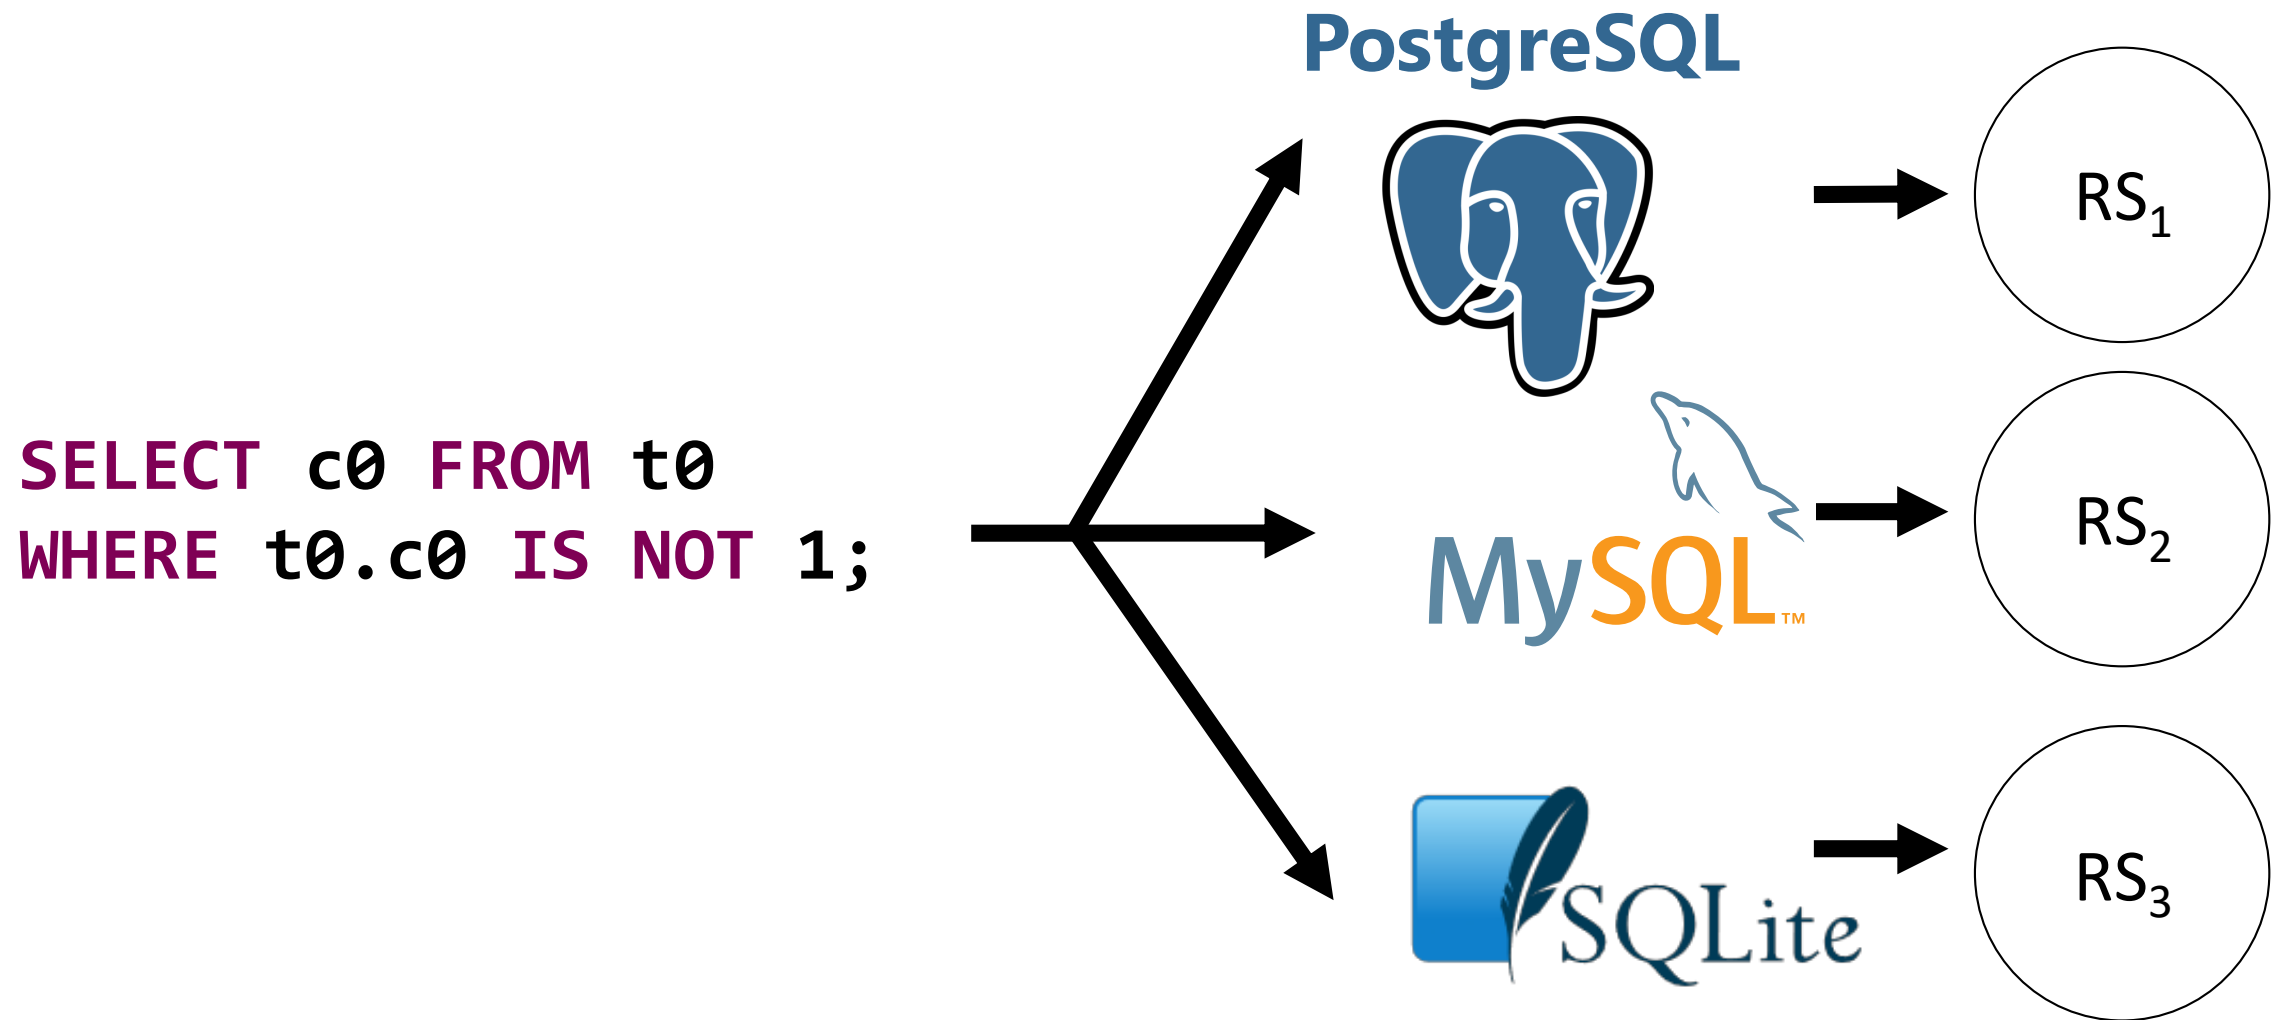
Background: Differential Testing

```
SELECT c0 FROM t0  
WHERE t0.c0 IS NOT 1;
```



Massive Stochastic Testing of SQL by Slutz, 1998.

Background: Differential Testing

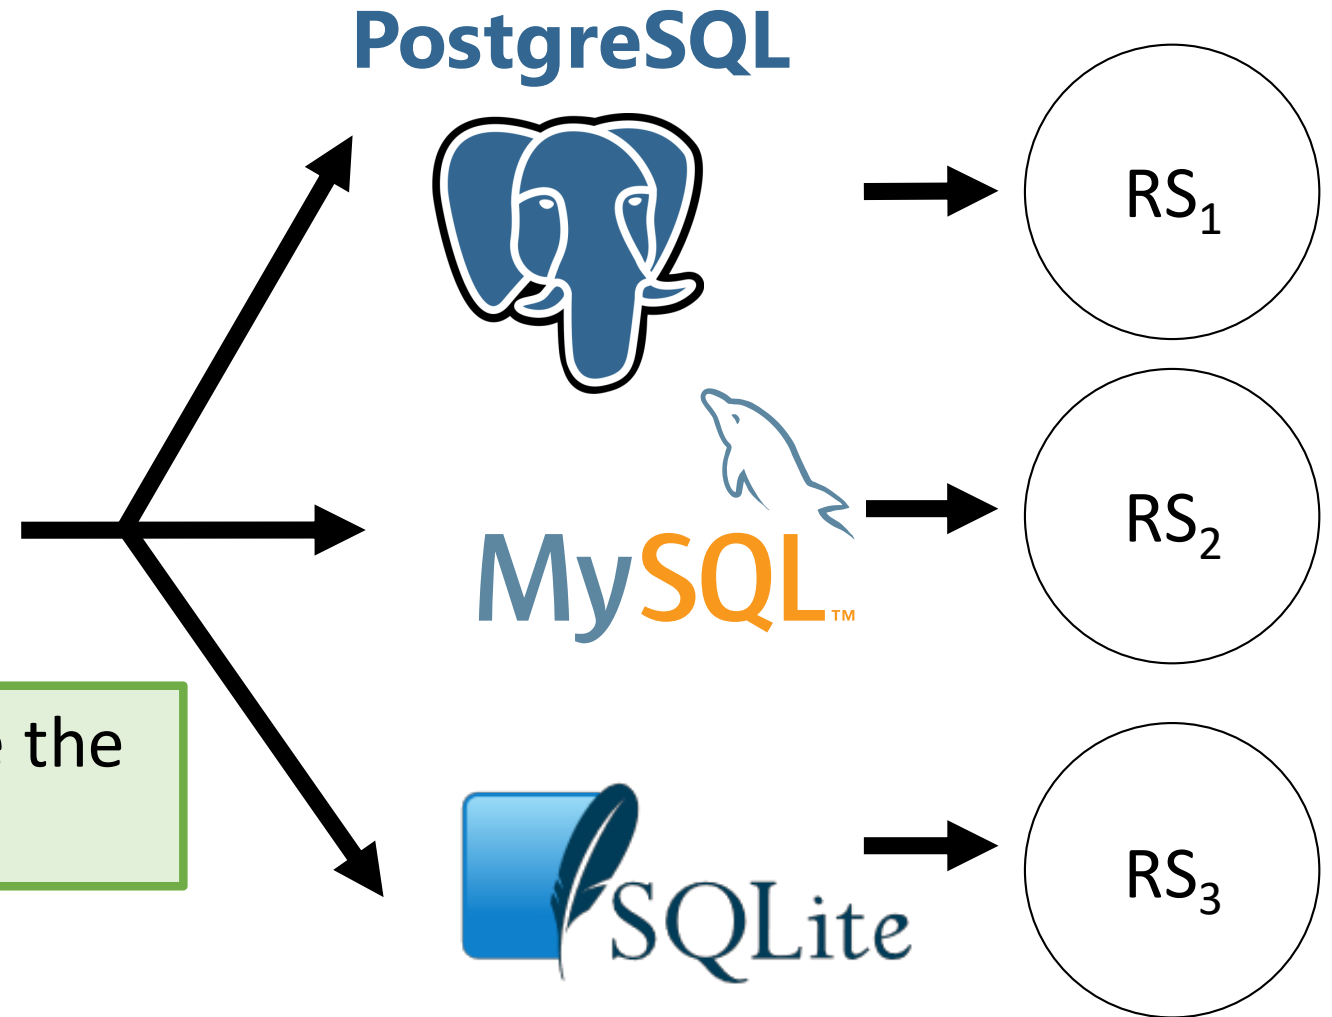


Massive Stochastic Testing of SQL by Slutz, 1998.

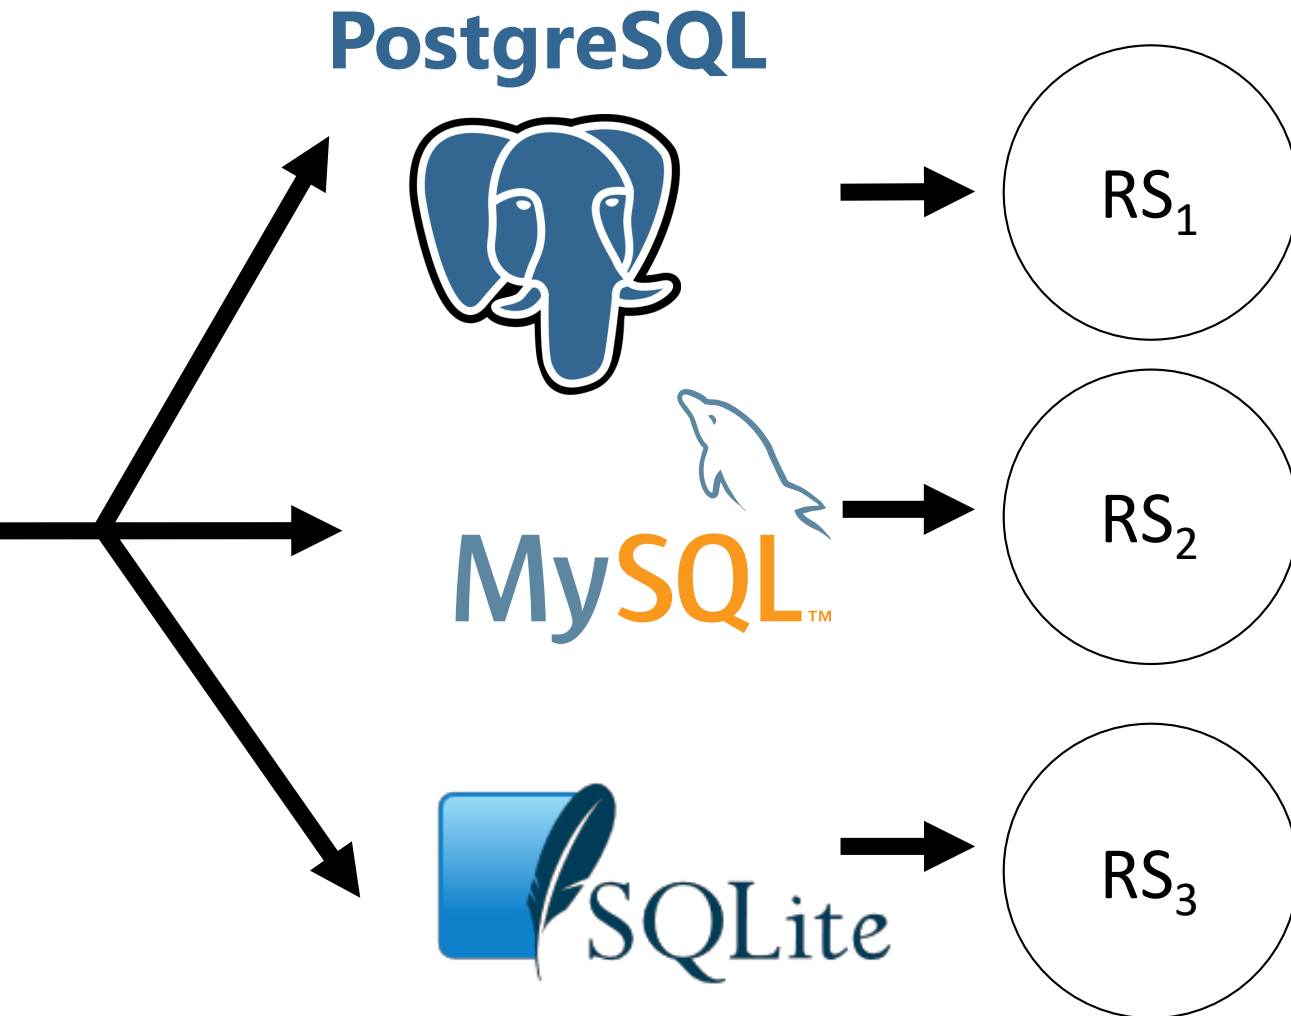
Background: Differential Testing

```
SELECT c0 FROM t0  
WHERE t0.c0 IS NOT 1;
```

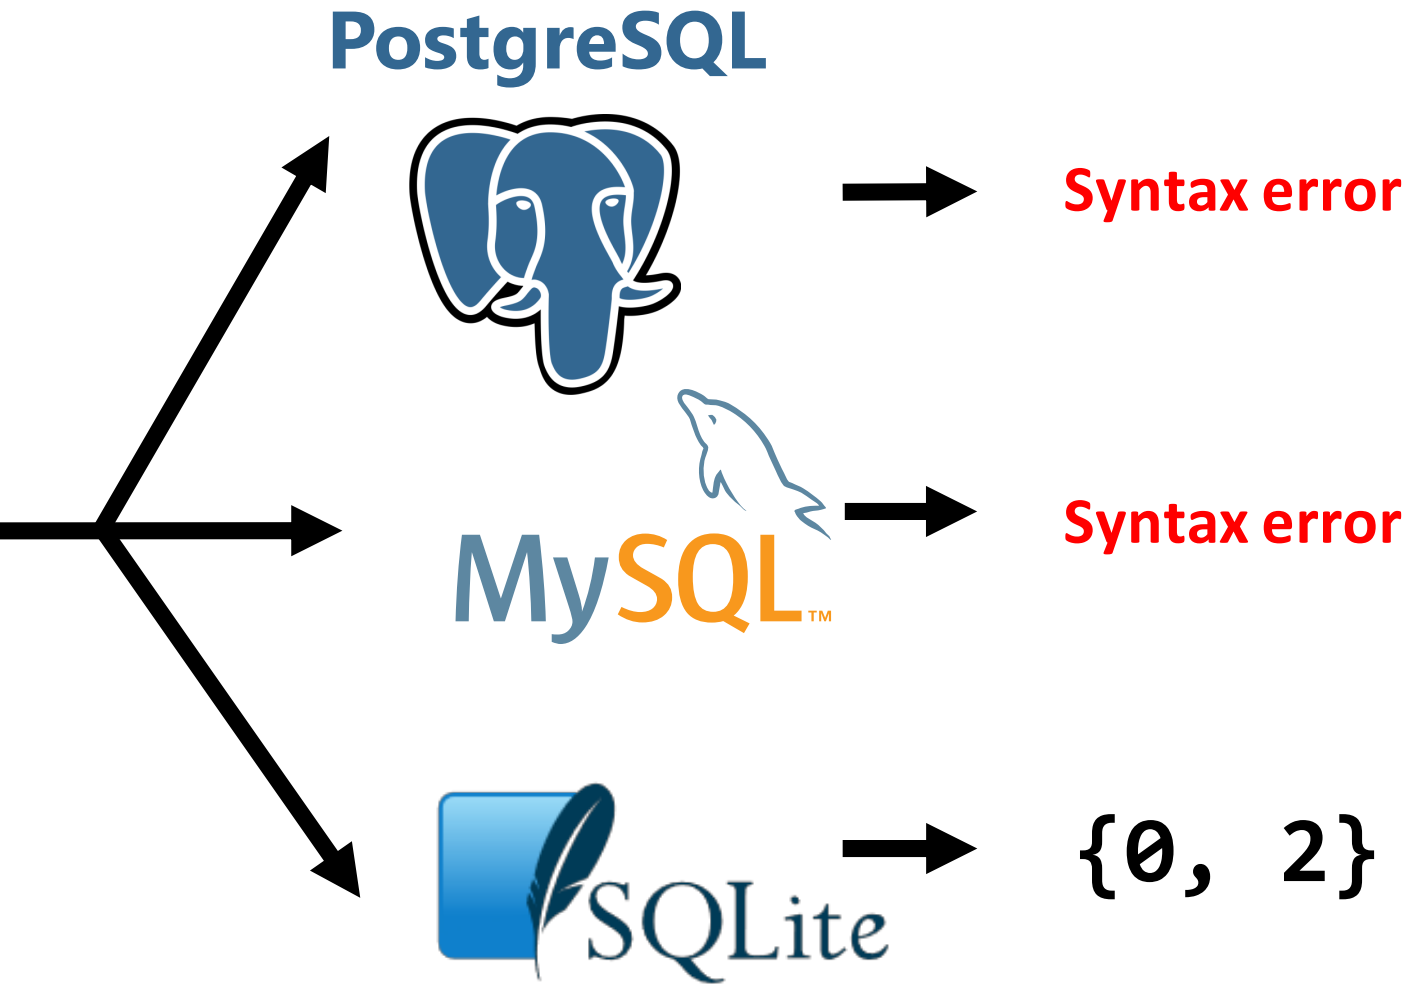
Check that all DBMSs compute the
same result ($RS_1 = RS_2 = RS_3$)



Background: Differential Testing



Background: Differential Testing



Background: Differential Testing

```
CREATE TABLE t0(c0);  
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;  
INSERT INTO t0 (c0) VALUES (0), (1), (2), (3), (NULL);  
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```

Background: Differential Testing

```
CREATE TABLE t0(c0);  
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;  
INSERT INTO t0 (c0) VALUES (0), (1), (2), (3), (NULL);  
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```

MySQL and PostgreSQL require a data type definition

Background: Differential Testing

```
CREATE TABLE t0(c0);  
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;  
INSERT INTO t0 (c0) VALUES (0), (1), (2), (3), (NULL);  
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```

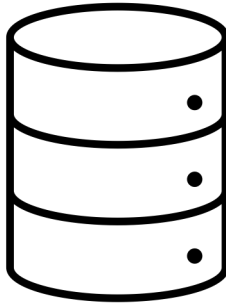
PostgreSQL provides an IS DISTINCT FROM operator,
and MySQL a <=> null-safe comparison operator

Pivoted Query Synthesis (PQS)

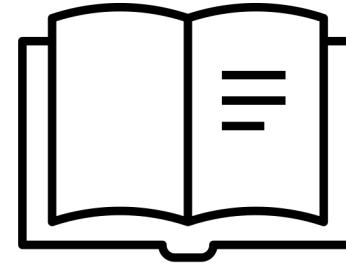
```
CREATE TABLE t0(c0);  
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;  
INSERT INTO t0 (c0) VALUES (0), (1), (2), (NULL);  
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```

t0

c0
0
1
2
NULL



0
2



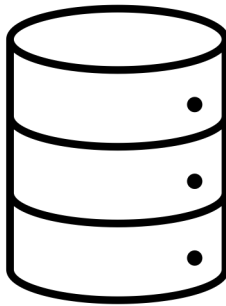
IS NOT is a “null-safe”...

Pivoted Query Synthesis (PQS)

```
CREATE TABLE t0(c0);  
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;  
INSERT INTO t0 (c0) VALUES (0), (1), (2), (NULL);  
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```

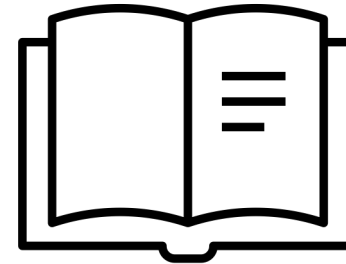
t0

c0
0
1
2
NULL



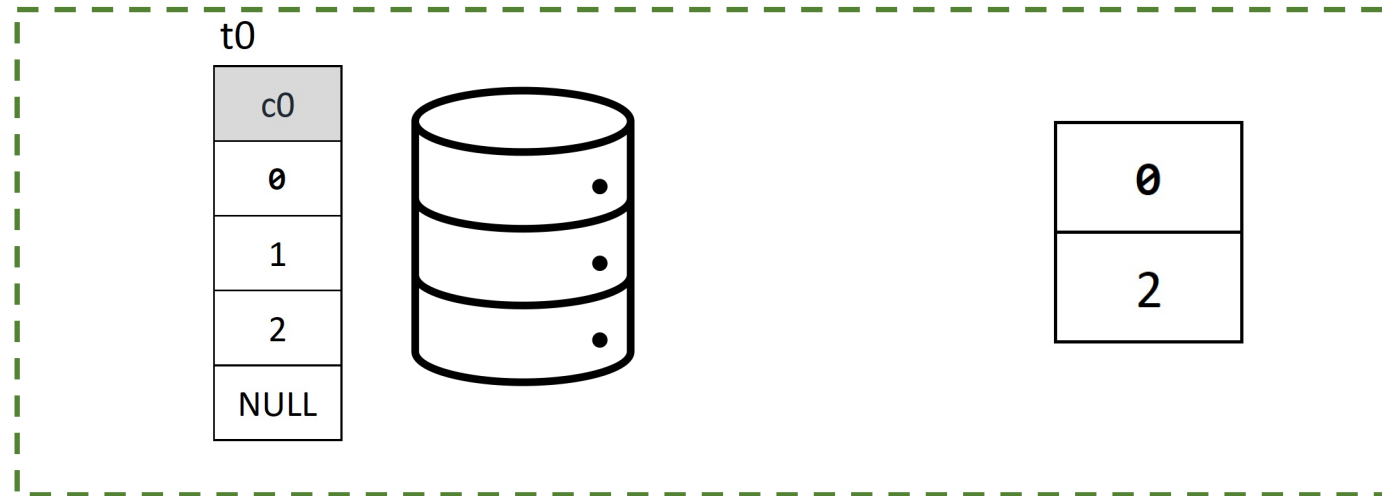
0
2

Correct or not?



IS NOT is a “null-safe”...

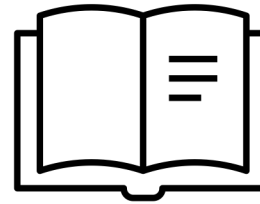
Pivoted Query Synthesis (PQS)



Pivoted Query Synthesis (PQS)

```
CREATE TABLE t0(c0);  
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;  
INSERT INTO t0 (c0) VALUES (0), (1), (2), (NULL);  
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```

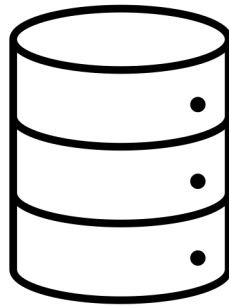
Query
Synthesis



IS NOT is a “null-safe”...

t0

c0
0
1
2
NULL



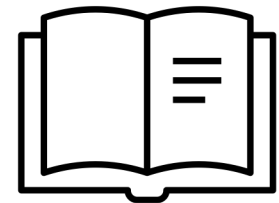
0
2

Pivoted Query Synthesis (PQS)

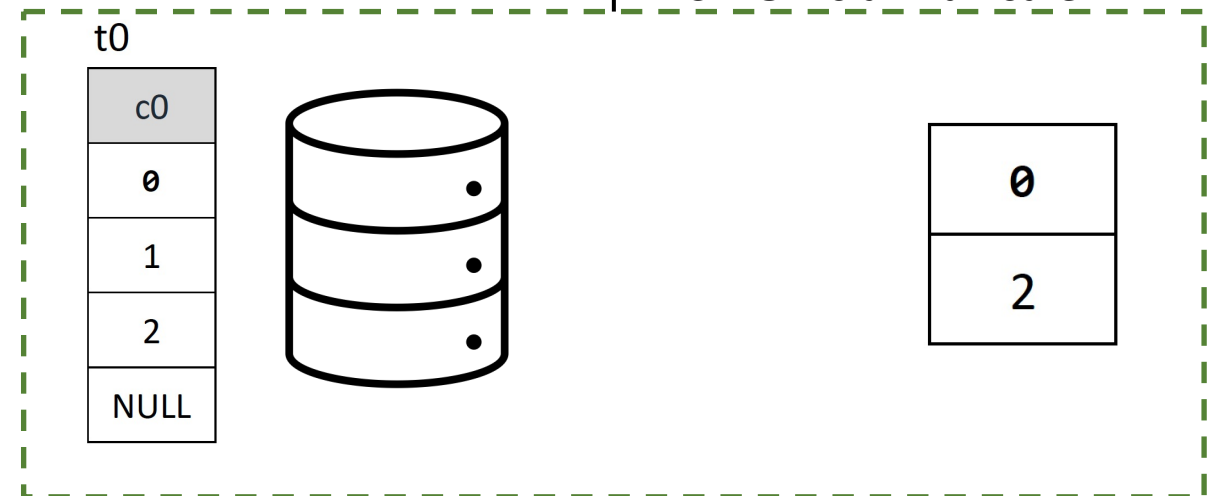
Infeasible to synthesize exact query for the output

```
CREATE TABLE t0(c0);  
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;  
INSERT INTO t0 (c0) VALUES (0), (1), (2), (NULL);  
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```

Query
Synthesis



IS NOT is a “null-safe” ...



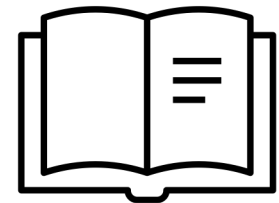
Pivoted Query Synthesis (PQS)

Infeasible to synthesize exact query for the output

Let's relax it: Synthesize a query producing partial output

```
CREATE TABLE t0(c0);  
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;  
INSERT INTO t0 (c0) VALUES (0), (1), (2), (NULL);  
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```

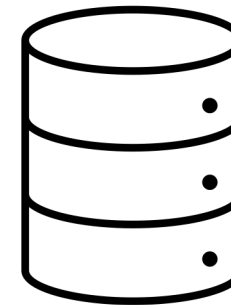
Query
Synthesis



IS NOT is a “null-safe” ...

t0

c0
0
1
2
NULL



0
2

PQS Idea

t0

c0
0
1
2
NULL

```
CREATE TABLE t0(c0);  
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;  
INSERT INTO t0 (c0) VALUES (0), (1), (2), (3), (NULL);  
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```

Validate the result set based on
one randomly-selected row

PQS Idea

t0

c0
0
1
2
NULL

```
CREATE TABLE t0(c0);  
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;  
INSERT INTO t0 (c0) VALUES (0), (1), (2), (3), (NULL);  
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```

← Pivot row

Validate the result set based on
one randomly-selected row

PQS Idea

t0

c0
0
1
2
NULL

```
CREATE TABLE t0(c0);  
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;  
INSERT INTO t0 (c0) VALUES (0), (1), (2), (3), (NULL);  
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```

NULL

TRUE

Generate a query that is
guaranteed to at least fetch
the pivot row

PQS Idea

t0

c0
0
1
2
NULL

```
CREATE TABLE t0(c0);  
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;  
INSERT INTO t0 (c0) VALUES (0), (1), (2), (3), (NULL);  
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```

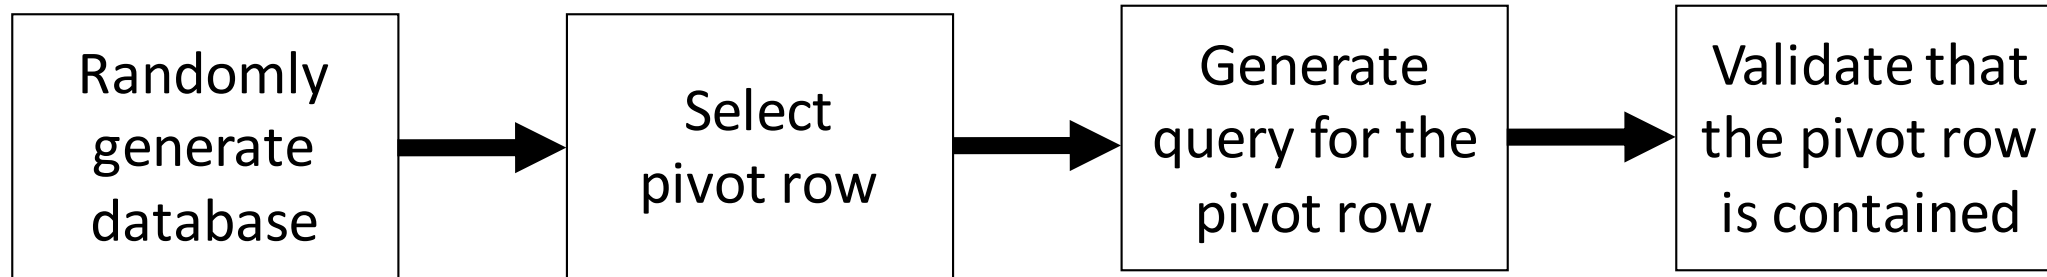


0
2

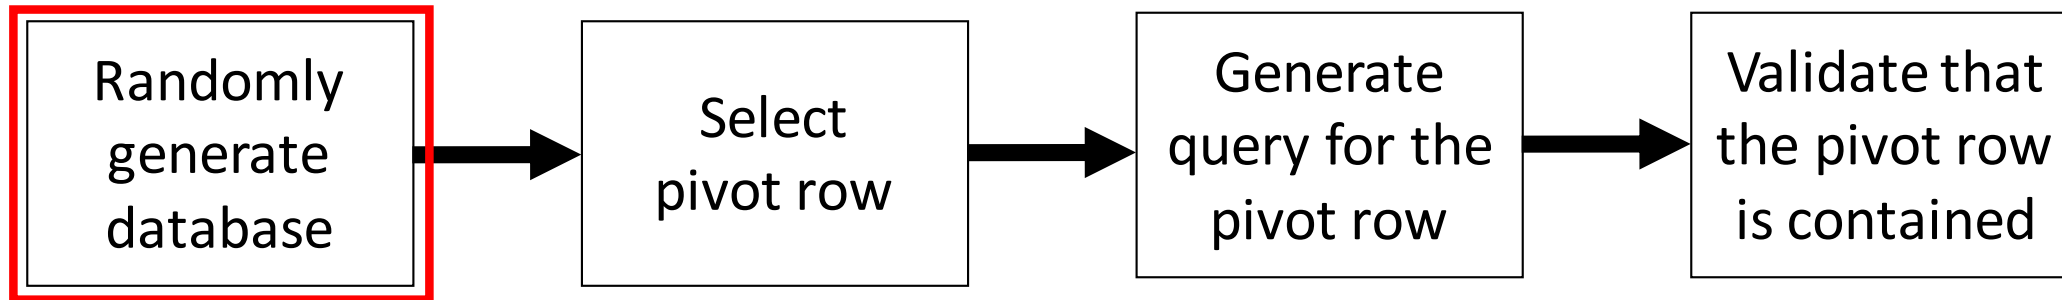


If the **pivot row** is missing from the result set a **bug** has been detected

Approach



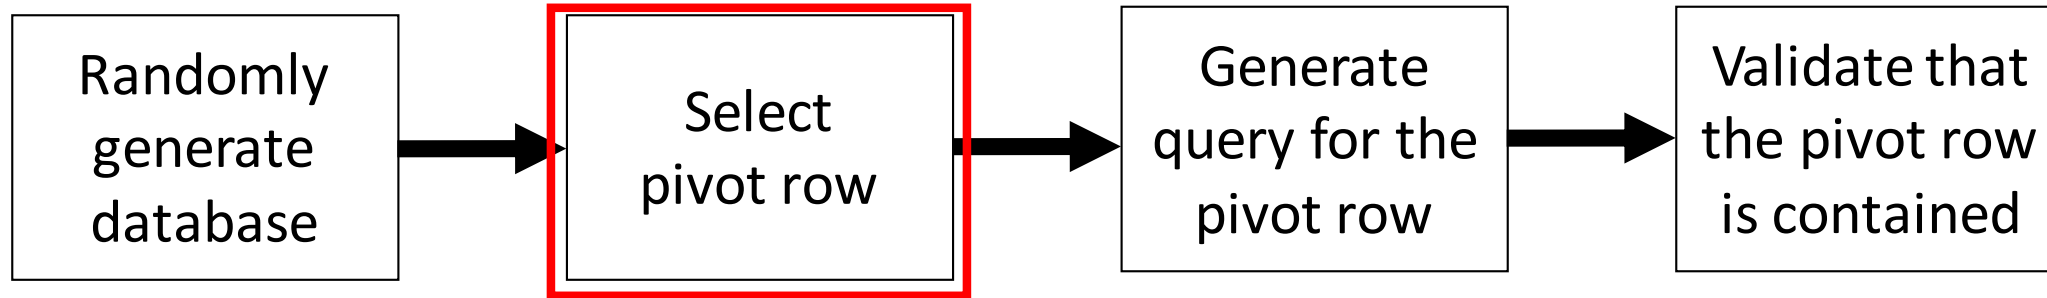
Approach



```
CREATE TABLE t0(c0);  
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;  
INSERT INTO t0 (c0) VALUES (0), (1), (2), (3), (NULL);
```

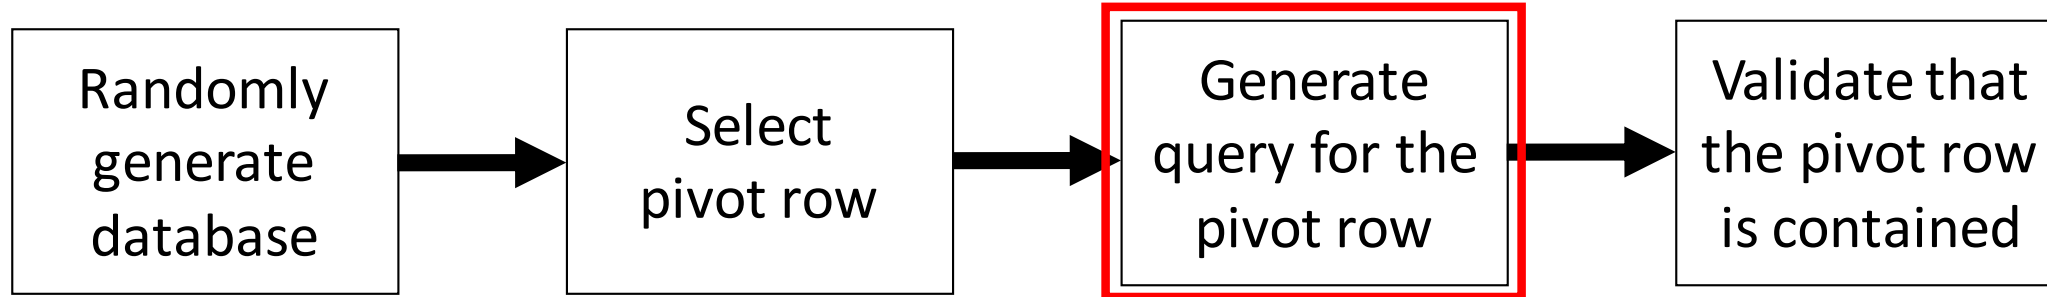
Statements are **heuristically generated**
based on the DBMS' SQL dialect

Approach



One **random row** from multiple tables and views

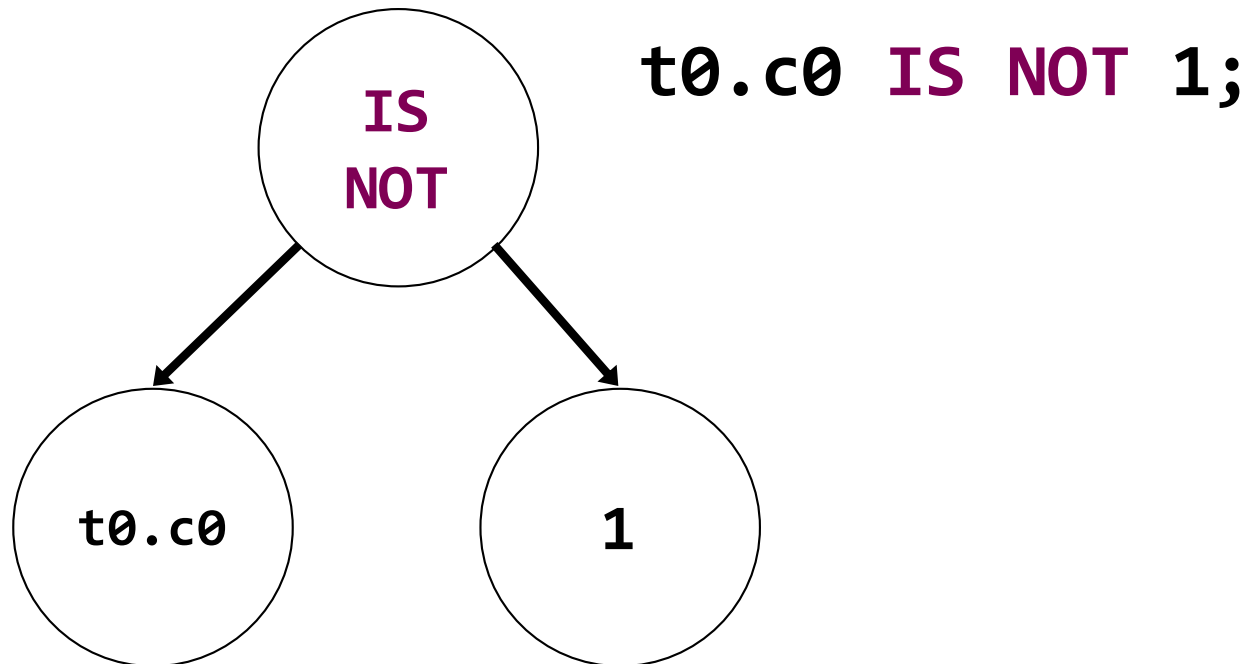
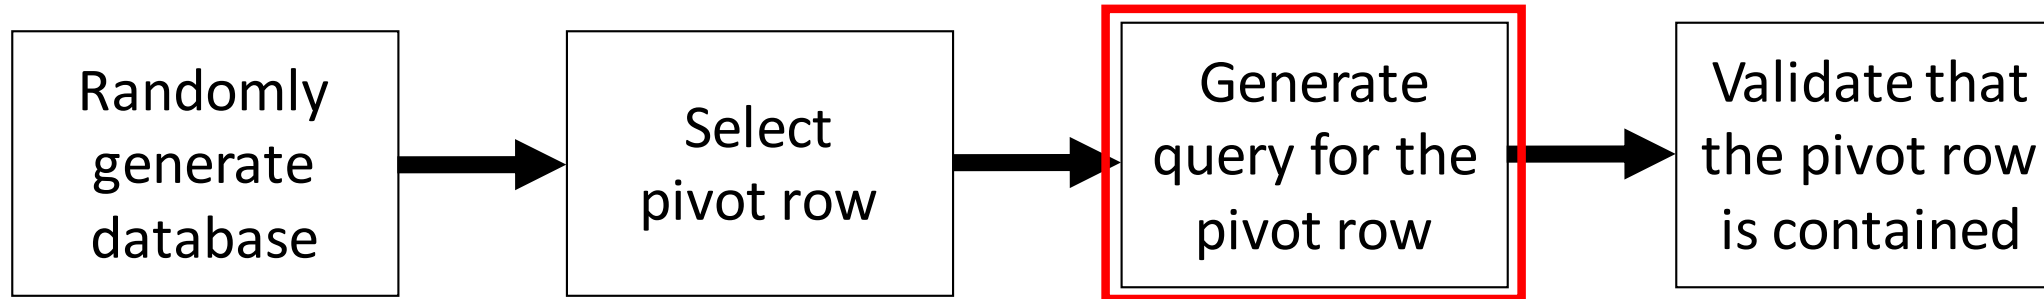
Approach



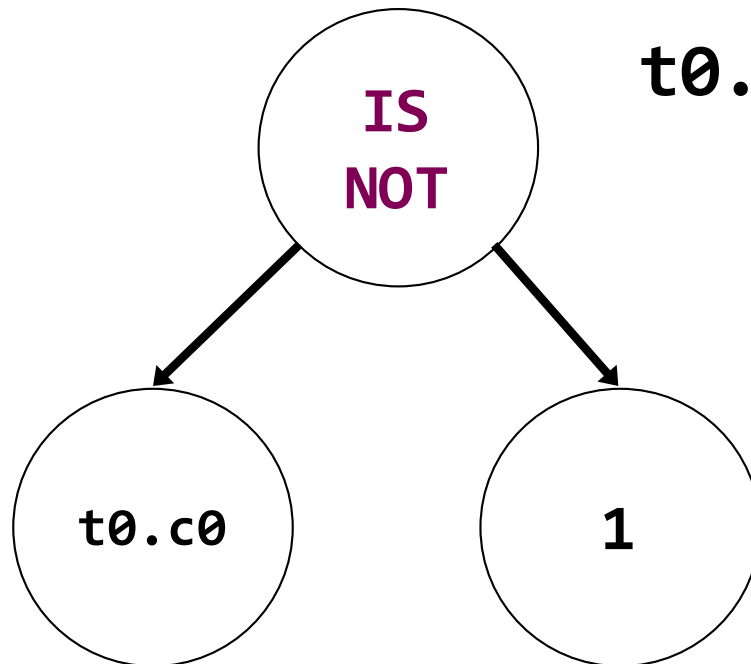
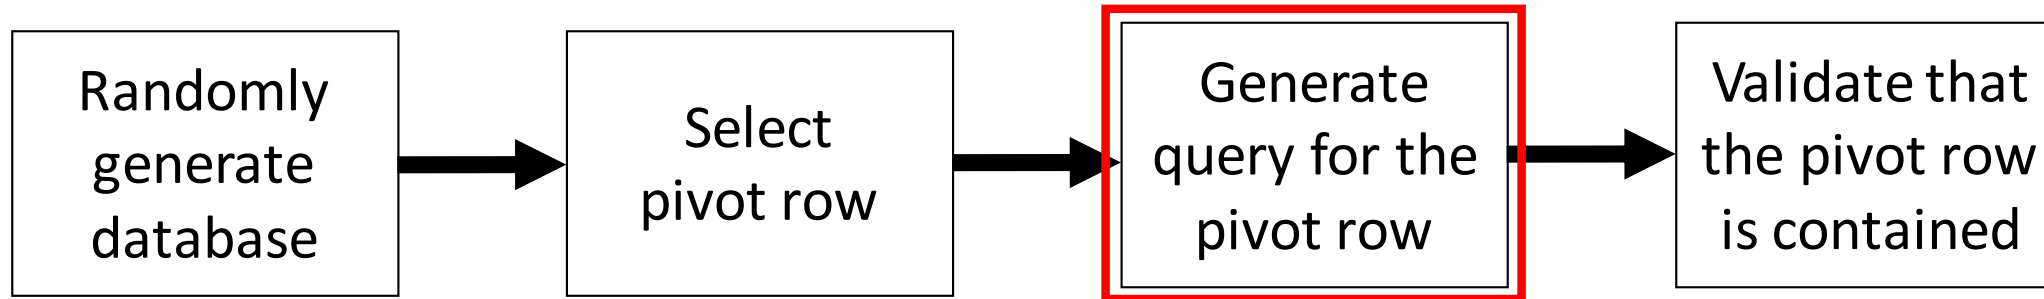
SELECT c0 **FROM** t0
WHERE

Generate **predicates** that **evaluate to TRUE** for the pivot row and use them in JOIN and WHERE clauses

Random Expression Generation



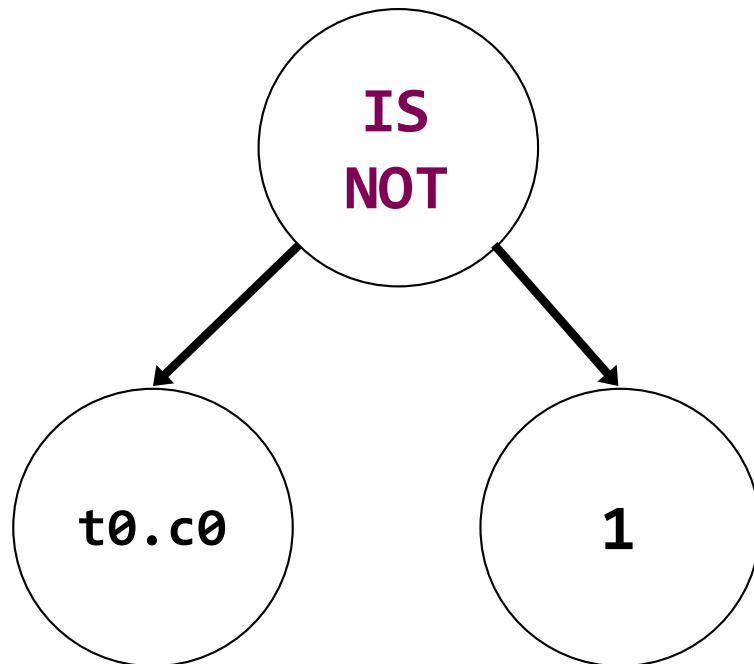
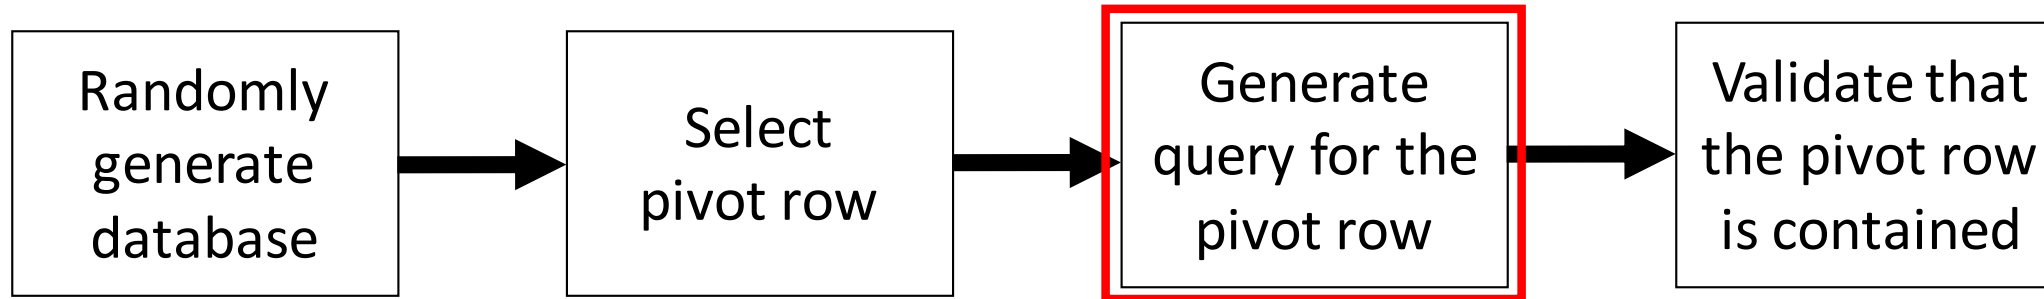
Random Expression Generation



`t0.c0 IS NOT 1;`

We implemented an **expression evaluator** for each node

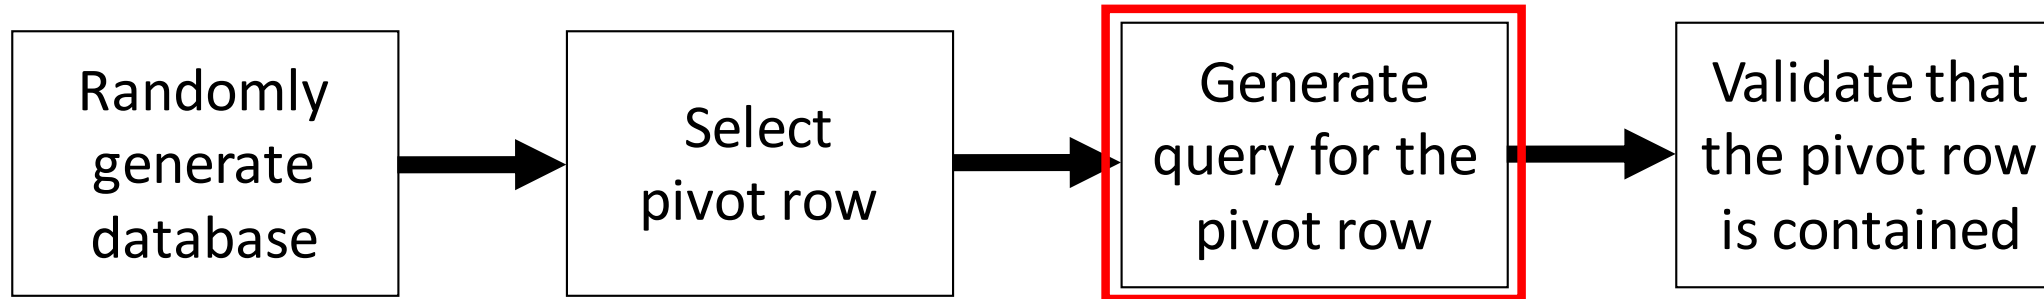
Random Expression Generation



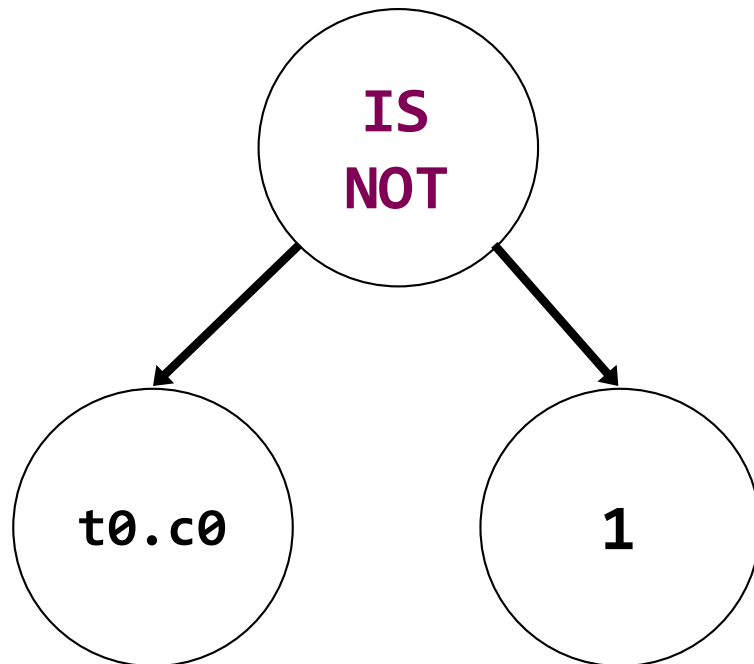
Evaluate the tree based on the **pivot row**

t0	
c0	
0	
1	
2	
NULL	

Random Expression Generation

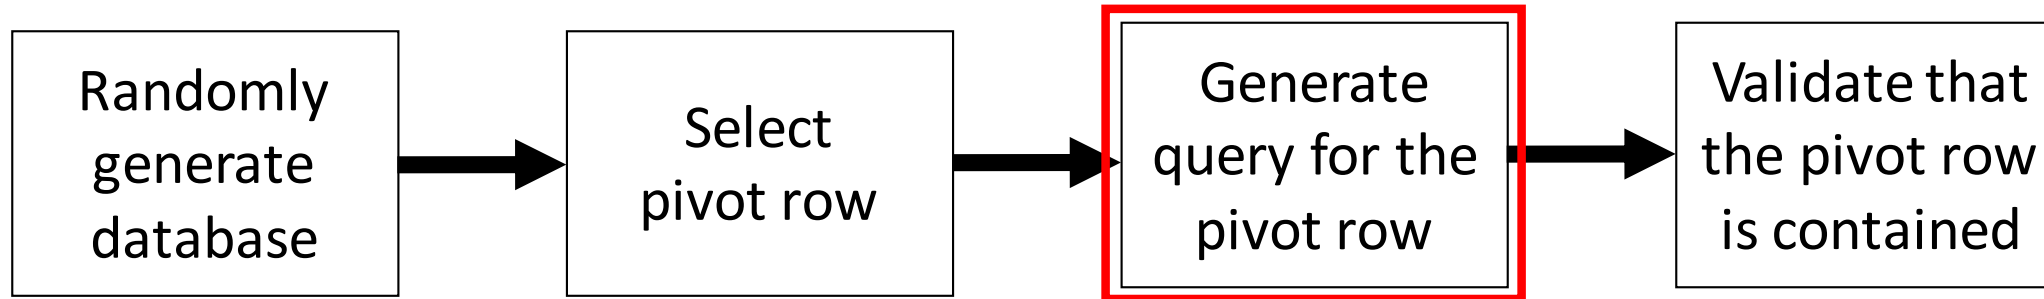


Column references return the values from the pivot row

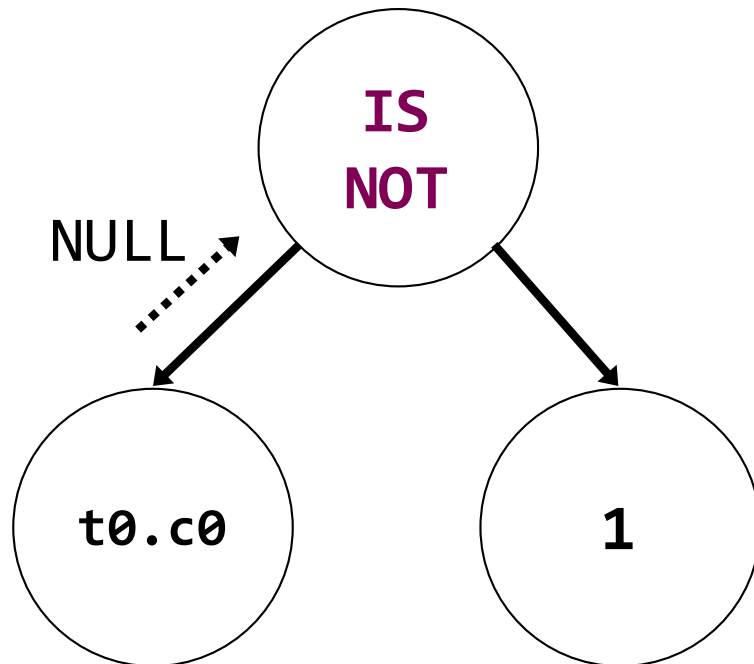


t0	
c0	
0	
1	
2	
NULL	

Random Expression Generation

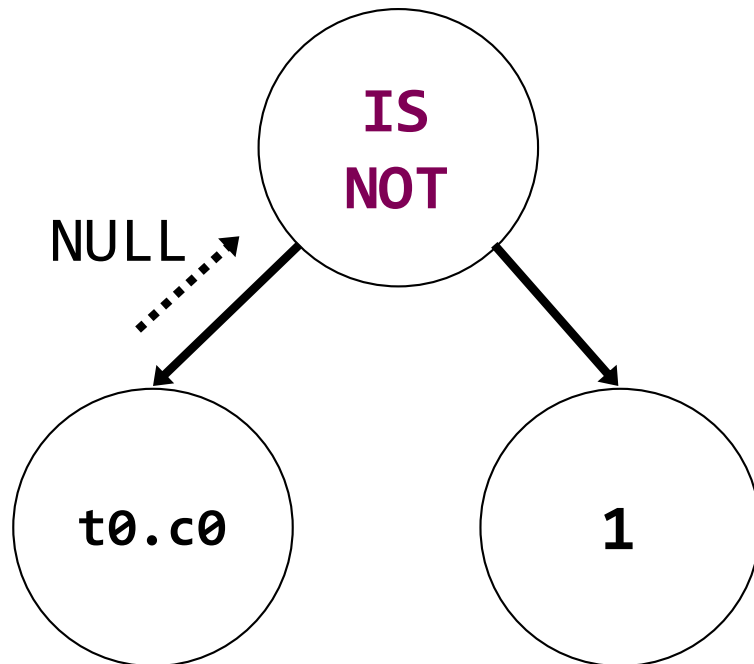
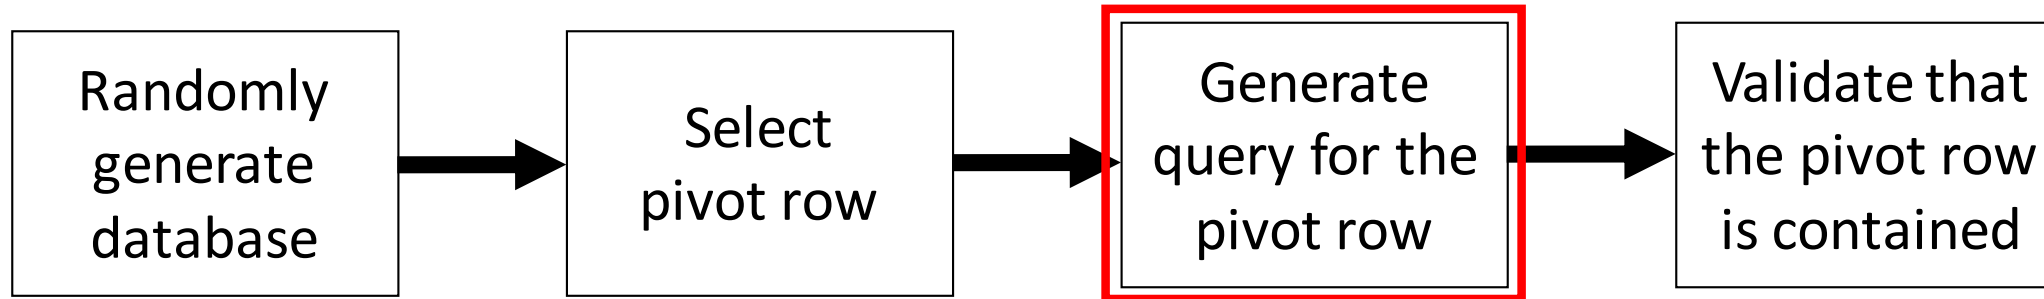


Column references return the values from the pivot row



t0
c0
0
1
2
NULL

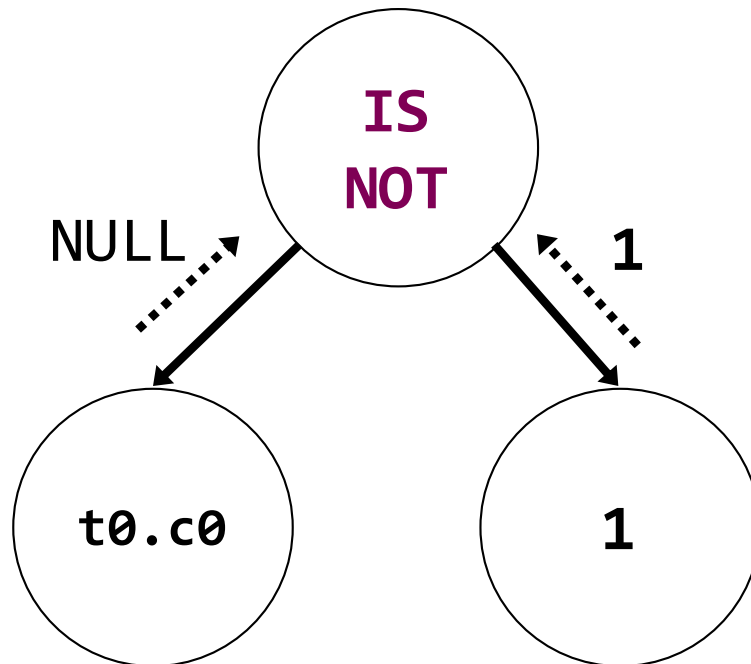
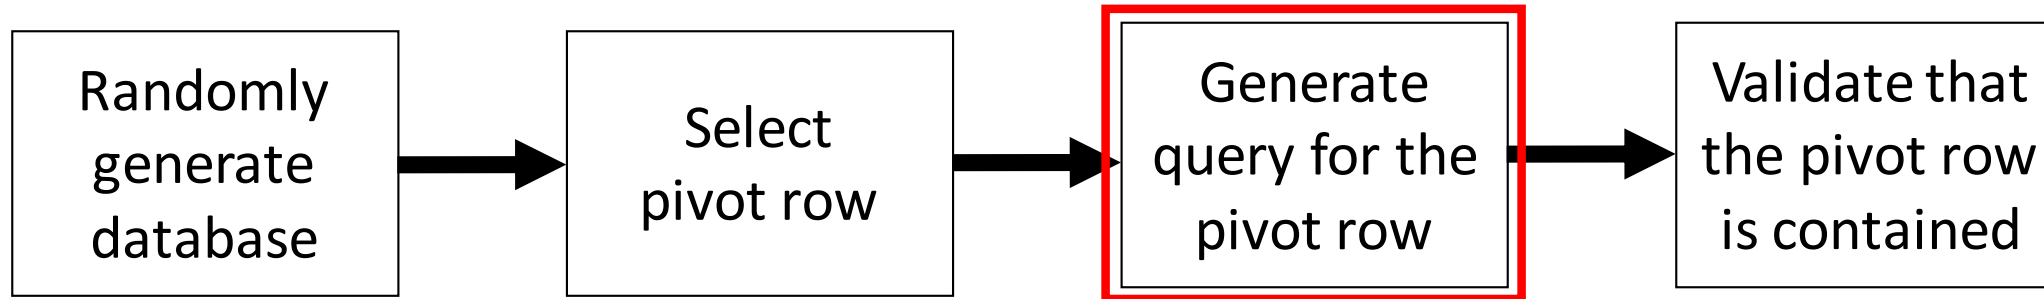
Random Expression Generation



Constant nodes return their assigned **literal values**

t0	
c0	
0	
1	
2	
NULL	

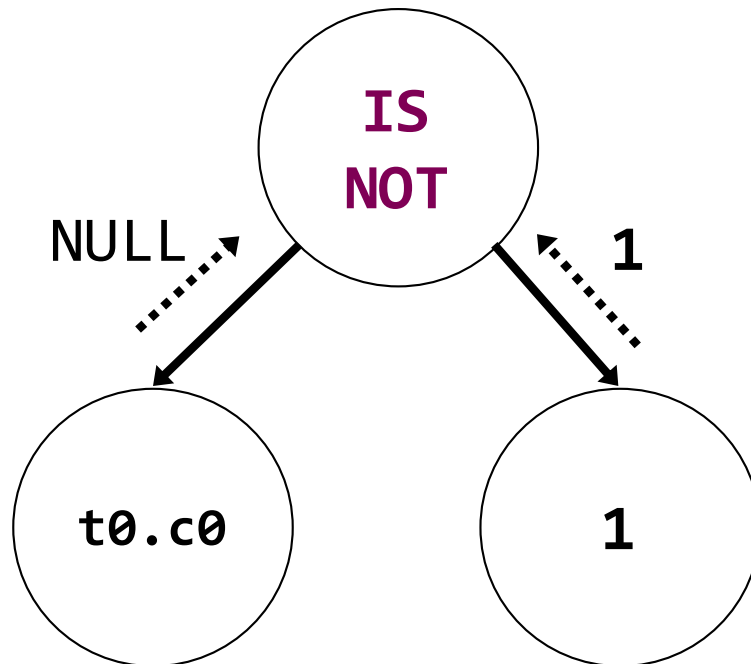
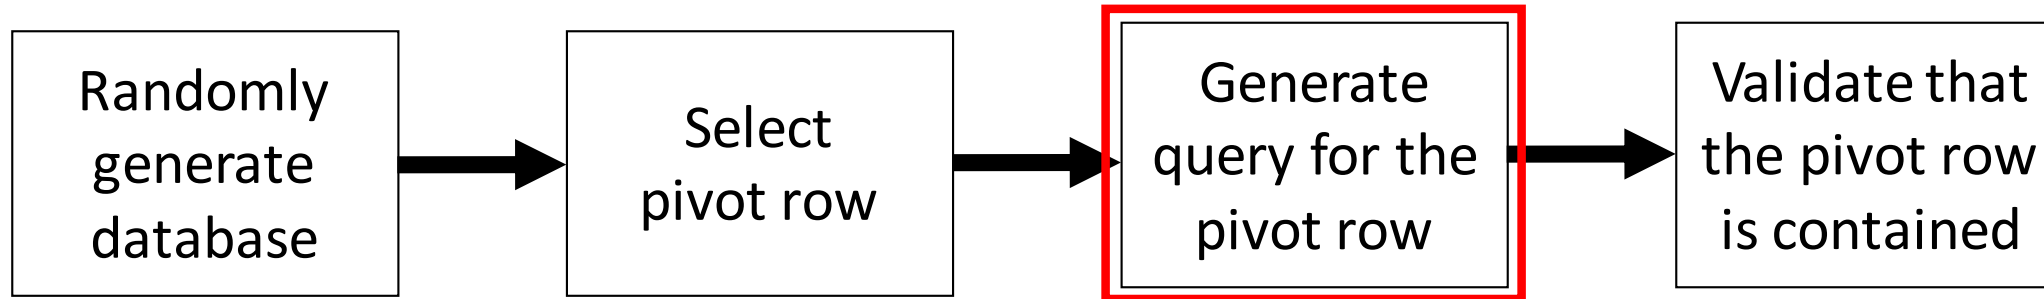
Random Expression Generation



Constant nodes return their assigned **literal values**

t0	
c0	
0	
1	
2	
NULL	

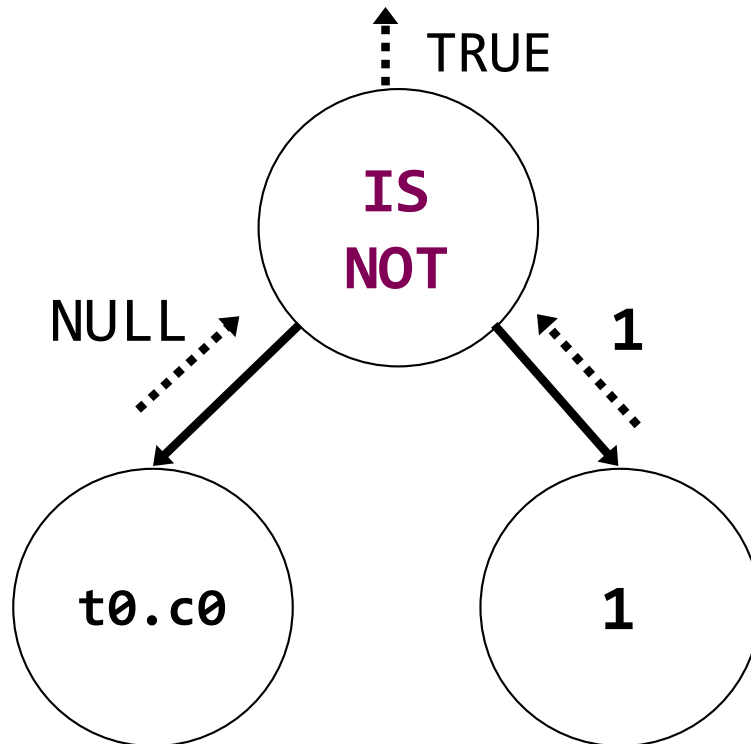
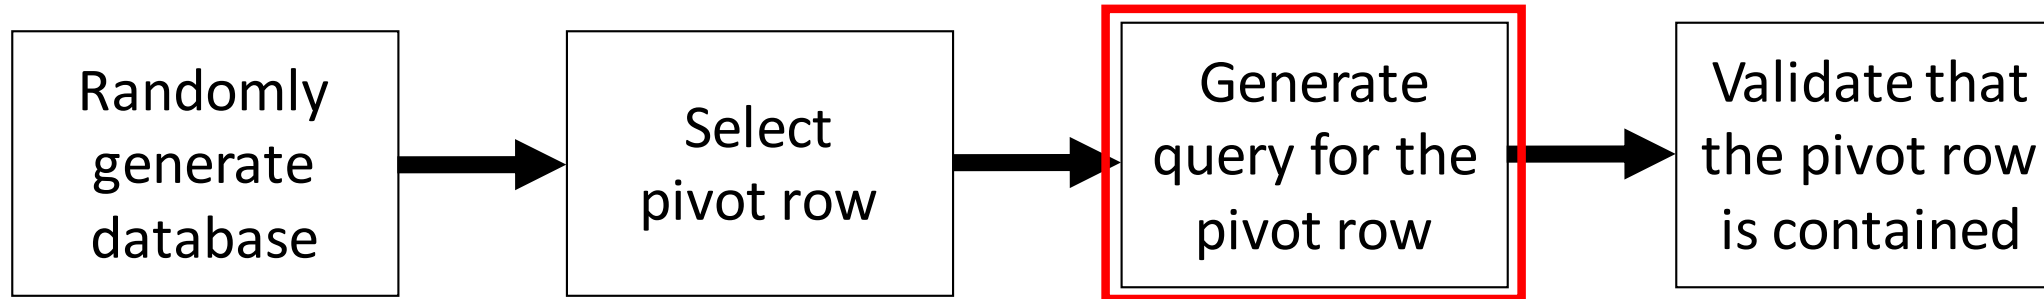
Random Expression Generation



Compound nodes
compute their result
based on their children

t0	
c0	
0	
1	
2	
NULL	

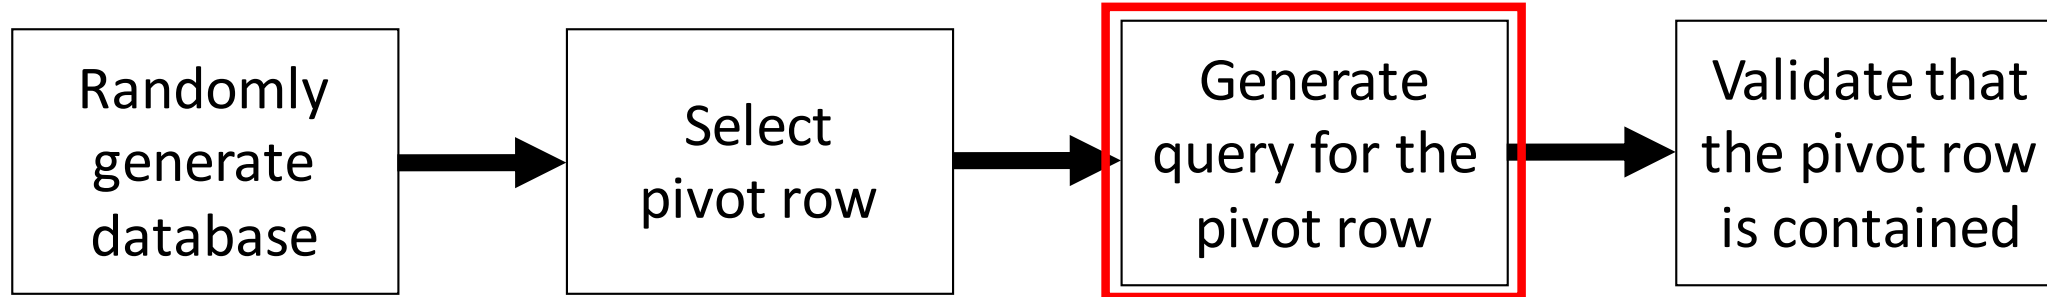
Random Expression Generation



Compound nodes
compute their result
based on their children

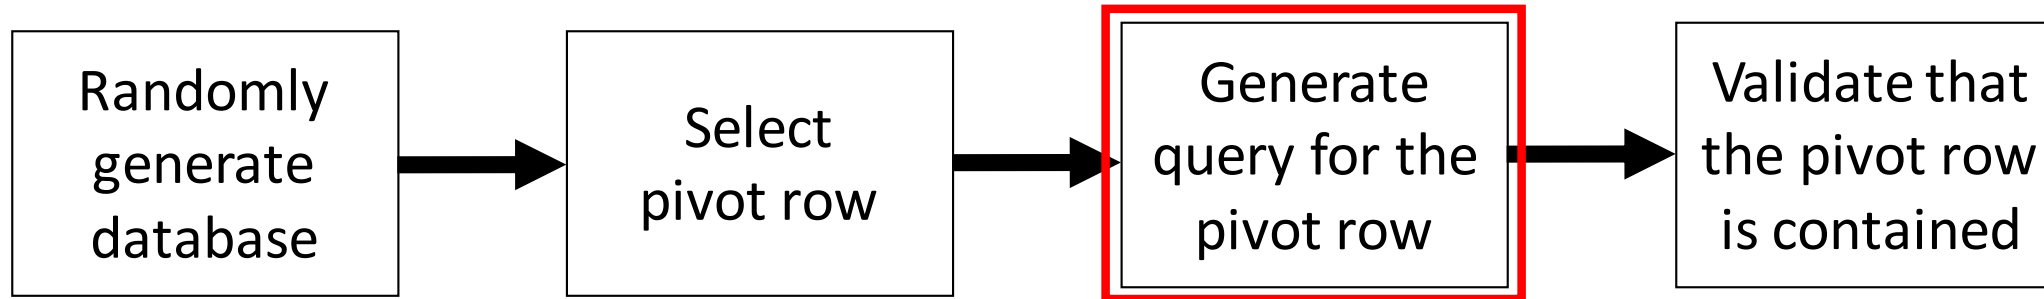
t0	
c0	
0	
1	
2	
NULL	

Query Synthesis



```
SELECT c0 c0 FROM t0
WHERE t0.c0 IS NOT 1;
```

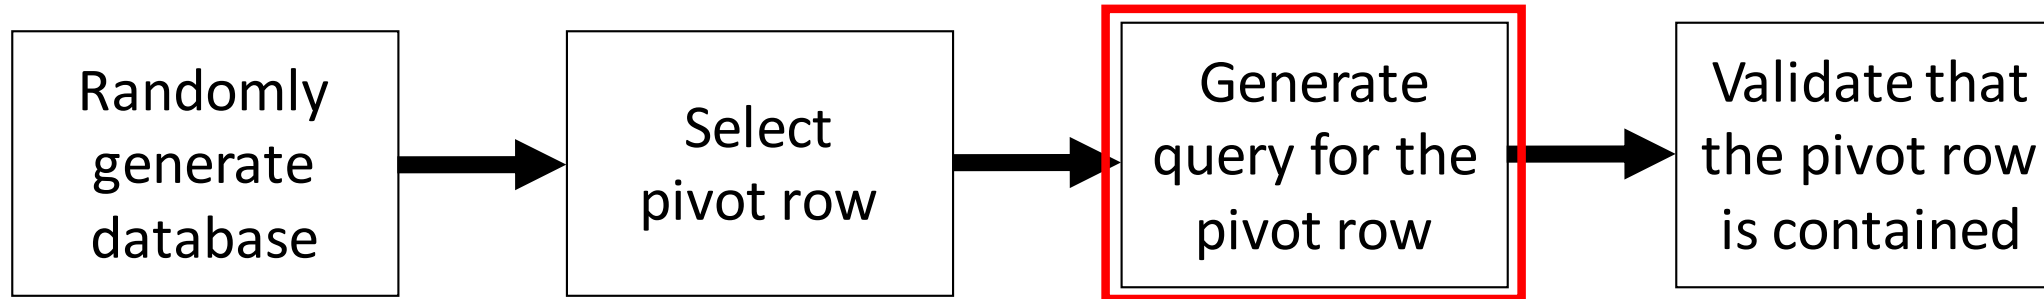
Query Synthesis



```
SELECT c0 c0 FROM t0  
WHERE t0.c0 IS NOT 1;
```

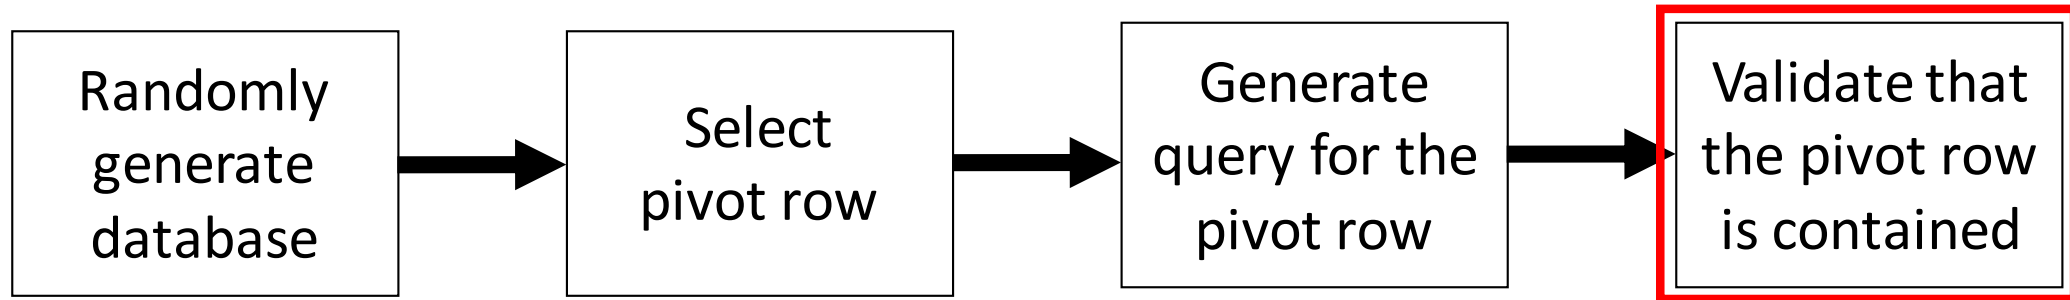
What if the expression **does not evaluate to TRUE?**

Random Expression Rectification



```
switch (result) {  
    case TRUE:  
        result = randexpr;  
    case FALSE:  
        result = NOT randexpr;  
    case NULL:  
        result = randexpr IS NULL;  
}
```

Approach

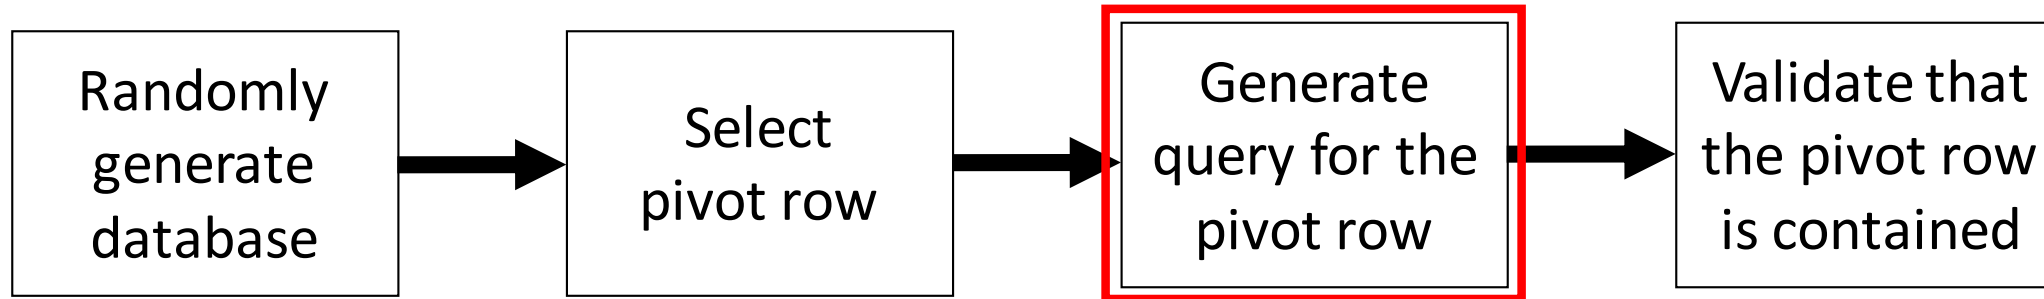


SELECT (NULL) INTERSECT

SELECT c0 FROM t0 WHERE NULL IS NOT 1;

Rely on the DBMS to check whether the row is contained

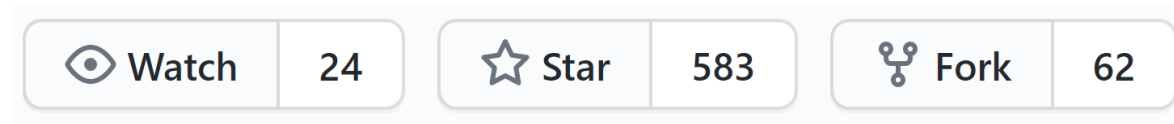
Random Expression Rectification



```
switch (result) {  
  case TRUE:  
    result = randexpr;  
  case FALSE:  
    result = NOT randexpr;  
  case NULL:  
    result = randexpr IS NULL;  
}
```

Alternatively, we could validate that the pivot row is expectedly **not fetched**

Implementation



<https://github.com/sqlancer>



Bugs Overview

DBMS	Fixed	Verified
SQLite	64	0
MySQL	17	7
PostgreSQL	5	3

Bugs Overview

DBMS	Fixed	Verified
SQLite	64	0
MySQL	17	7
PostgreSQL	5	3

96 bugs were unique,
previously unknown ones

Oracles

DBMS	Logic	Error	Crash
SQLite	46	17	2
MySQL	14	10	1
PostgreSQL	1	7	1

61 were logic bugs

Discussion: Bug Importance

```
CREATE TABLE t0 (c0);  
CREATE TABLE t1 (c1);  
INSERT INTO t0 VALUES (1);  
SELECT c0 FROM t0 LEFT JOIN t1 ON c1=c0 WHERE NOT (c1 IS NOT NULL AND c1=2);
```


Discussion: Bug Importance

```
CREATE TABLE t0 (c0);  
CREATE TABLE t1 (c1);  
INSERT INTO t0 VALUES (1);  
SELECT c0 FROM t0 LEFT JOIN t1 ON c1=c0 WHERE NOT (c1 IS NOT NULL AND c1=2);
```

This is a cut-down example, right? You can't possibly mean to do that WHERE clause in production code.

Discussion: Bug Importance

```
CREATE TABLE t0 (c0);  
CREATE TABLE t1 (c1);  
INSERT INTO t0 VALUES (1);  
SELECT c0 FROM t0 LEFT JOIN t1 ON c1=c0 WHERE NOT (c1 IS NOT NULL AND c1=2);
```

This is a cut-down example, right? You can't possibly mean to do that WHERE clause in production code.

*I might not spell it like that myself, but a **code generator** would do it (and much worse!). This example was **simplified from a query generated by a Django ORM** queryset using `.exclude(nullable_joined_table__column=1)`, for instance.*

Discussion: Bug Importance

```
CREATE TABLE t0 (c0);  
CREATE TABLE t1 (c1);  
INSERT INTO t0 VALUES (1);  
SELECT c0 FROM t0 LEFT JOIN t1 ON c1=c0 WHERE NOT (c1 IS NOT NULL AND c1=2);
```

This is a cut-down example, right? You can't possibly mean to do that WHERE clause in production code.

*I might not spell it like that myself, but a **code generator** would do it (and much worse!). This example was **simplified from a query generated by a Django ORM** queryset using `.exclude(nullable_joined_table__column=1)`, for instance.*

Even “obscure” bugs might affect users

Discussion: Limitations

- Implementation effort for complex operations
- Requires understanding of the SQL semantics
- Aggregate and window functions
- Ordering
- Duplicate rows

Discussion

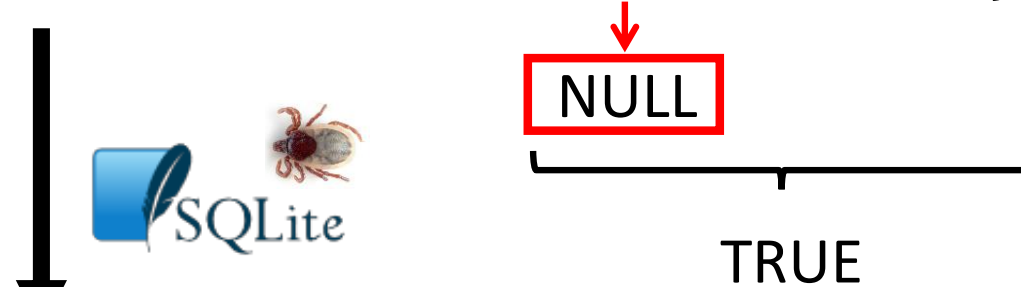
- Does not explain the bugs well
 - Nature of black-box testing
- SQLancer detects bugs that expected results not outputted (containment bugs?)
 - Unexpected results can rarely be outputted
 - Ensure the completeness of results? Hard!
- SQLancer mostly tests the boolearn evaluator within DBMS
 - The WHERE and JOIN clauses are evaluated correctly or not
 - Reminds me of YinYang paper (PLDI '20)

Example: SQLite3 Bug

t0

c0
0
1
2
NULL

```
CREATE TABLE t0(c0);  
CREATE INDEX i0 ON t0(1) WHERE c0 NOT NULL;  
INSERT INTO t0 (c0) VALUES (0), (1), (2), (NULL);  
SELECT c0 FROM t0 WHERE t0.c0 IS NOT 1;
```



0
2

NULL was not contained
in the result set!

